

Mynyra Project Lifecycle and Research Workflow

The strongest supported gap is the operational connection between existing controls, project status and repeated work. Mynyra already has phases, decision gates, scripts and a relational research catalog. What remains unproven is whether these pieces let the people operating the project start, finish and hand over each research run with progressively less coordination effort.

The reported symptom is specific: two research campaigns felt similarly difficult to execute and coordinate. That deserves investigation in its own right. It cannot be explained away by calling a rejected strategy an informative result, and it does not by itself establish that a different database would remove the friction.

Finding	Assessment
Existing technical structure	Confirmed in the inspected repository: an economic worldmap, implementation increments, four research work packages, fixed experiment definitions, SQLite metadata and bounded scripts. ^{1,3,8}
Weak shared operational visibility	Supported concern: top-level summaries lag the expanded result, while detailed evidence is distributed among newer records. ^{2,7}
Repeatable workflow integration	Supported concern: successive work packages encountered identity/integration defects, and the current runner binds to a specific earlier proof directory. ^{5,6,9}
Exact cause and cost of repeated friction	Unmeasured: no operator-time breakdown or observation of both workstations establishes the dominant source of delay.

A suitable response is a **hybrid lifecycle with nine phases**, a bounded research subproject where justified, and a small shared status-and-handoff mechanism that reuses the current evidence system. This is a proposed operating design, not an approved change to the repository or its gates.

The economic target remains repeatable, actually received cash that materially finances continued operation. The latest inspected strategy state is no advancement; January-March is exploratory, April remains sealed, and no shadow campaign or order experiment follows automatically.^{7,26}

The economic urgency described by the project participants makes operator time and avoidable delay material costs. The exchange-rate and consumer-price examples are treated as reported context, not independently verified market statistics. No conclusion depends on their precise amounts.

1. py-mynyra, WORLDMAP.md.
3. py-mynyra, Locked research expansion sequence.
7. py-mynyra, Praxis Step 4: expanded exploratory comparison.

Repository evidence and the management gap

The inspection covered the complete, non-truncated tree of 72 tracked files, including 31 documents, at main commit fd016456205e3ef35c9b7ec487fad51362c516e4. Key plans, protocols, catalog schema and runner code were read, alongside the available PR #1-#12 metadata and descriptions. This is a static assessment; private numerical evidence and runtime services were not independently replayed.

Evidence	Practical implication
WORLDMAP and implementation plan	Already connect technical capabilities to account access, receipts, repeatability and open economic decisions. "No lifecycle throughout the repository" is not supported. ^{1,2}
Locked expansion sequence	Defines four work packages and completion dependencies; its closeout explicitly creates no automatic fifth step. ³
SQLite migration and Step 3 proof	Already model datasets, experiments, snapshots, runs and artifact relationships. The report records baseline reconciliation and isolated restore; that establishes documented local capability, not a shared remote service. ^{5,8}
Top-level status versus Step 4	README and implementation summaries still foreground the initial five-candidate result while newer records close the expanded batch. A reader must reconcile several levels of status. ^{2,7}
Identity repairs across packages	Step 3 reports an omitted chronological run-key dimension. Step 4 reports a missing experiment ID in result envelopes and 523 retained partial artifacts. These are concrete interface-learning costs. ^{5,6}

The catalog schema tracks research evidence rather than the full operating workflow: it has no dedicated phase, task, next-action, responsible-person or approval-event relations. That is a scope observation, not a demand to add every concept to SQL. Documents or existing task records can remain authoritative for these facts.⁸

The Step 4 runner restores a named Step 3 catalog and uses a named snapshot path. This is defensible for a frozen historical experiment. It also shows why a future reusable interface should live alongside frozen runners rather than rewrite their provenance.⁹

A related Mynyra-Trade charter already describes daily review, evidence and automation, but applies to an older engine/Control Room context. Its champion, database-authority and seven-day demo requirements are not silently imported into py-mynyra.²⁷

2. py-mynyra, IMPLEMENTATION_PLAN.md and README.md.

5. py-mynyra, Praxis Step 3: one frozen data-to-result path.

6. py-mynyra, Step 4 artifact identity correction.

8. py-mynyra, Authoritative local SQLite catalog migration.

9. py-mynyra, Praxis and bounded Step 4 runner.

Terminology and research basis

Lifecycle means the overall progression of a project. Its parts are **phases**. A phase delivers something needed for a decision. A workflow describes how activities and handoffs happen; a gate is the decision about whether to proceed. PMI explicitly distinguishes project phases from its management process groups.¹⁰

Work breakdown structure (WBS) divides scope into deliverable-oriented pieces; a work package is a manageable element at the bottom of that decomposition. A large phase can contain subphases or several work packages. It deserves a separate subproject only when a distinct outcome, accountable lead, schedule, resources and acceptance decision make separate management useful.²⁵

RDBMS means relational database management system; SQLite already qualifies. **MCP** means Model Context Protocol. It standardizes how an AI application accesses tools and context. The specification does not determine how an application manages the context it receives. A shared data model, permissions, freshness and workflow rules still have to be defined.²⁰

Research basis	Reason for using it here
PMI and APM lifecycle guidance	Supports explicit deliverables and phase decisions, tailored to uncertainty; hybrid structure is appropriate when research iterates and later exposure decisions need firm boundaries. ^{10,11}
APM benefits management	Keeps delivered software and realized business benefit separately measurable; the income objective needs receipts and operating sustainability. ¹²
IBM CRISP-DM	Provides a recognizable research progression from business and data understanding to modeling, evaluation and use of results. ^{13,28}
2026 model-risk guidance and FCA review	Support purpose-bounded models, accountable review, testing and ongoing monitoring in financial work. ^{14,15,16}
BCBS data principles and DORA flow guidance	Inform authoritative definitions, lineage, timely reporting and inspection of handoff bottlenecks. ^{17,19}
Google SRE	Separates repetitive automatable operational work from necessary judgment and lasting engineering improvements. ¹⁸

These sources do not prescribe a universal nine-phase financial-project standard. The phases below are an explicit synthesis for Mynyra. Banking supervision and the FCA review provide transferable practices, not a finding that those regimes apply to this Iranian team. SR 26-2 replaced SR 11-7 in April 2026 and stresses proportionality; a two-person research effort should not copy a large bank’s organization chart.¹⁴

Ongoing income generation is an operation after establishment. Multiple related research and implementation projects may support it. Keeping those boundaries visible prevents an indefinitely expanding software project from becoming the business objective.

10. Project Management Institute, Projects and the Project Lifecycle.
11. Association for Project Management, What is a life cycle?.
12. Association for Project Management, What is benefits management and project success?.
14. Federal Reserve, OCC and FDIC, SR 26-2: Revised Guidance on Model Risk Management.
20. Model Context Protocol, Architecture overview.
25. Association for Project Management, Project management glossary.

Phases 1-3: decide what work deserves time

1. Economic baseline and project definition

Why chosen: a cash-constrained project needs a measurable destination and a current limit on what it can afford to consume.

Confirm the private baseline for usable cash, recurring costs, reserve needs, available operator hours and survivable losses. Define what counts as self-sustaining income, including its measurement window and actual cash receipts. Reuse the canonical objective and worldmap; supply missing values rather than rewrite their prose.¹

Output and exit: a short current baseline with dated observations, owner, unresolved facts and approved time/spend bounds for the next work package. Unknown personal limits do not block ordinary read-only research, but consequential actions cannot assume them. Revisit this phase when circumstances materially change.

2. Route feasibility and comparative business case

Why chosen: it prevents a technically interesting route from consuming scarce time before its practical ability to support the income objective is examined.

Compare a small number of admissible routes against time to useful evidence, required capital, data access, operational burden, recovery and cash-receipt feasibility. An automated trading route is a candidate; manual procedures or another justified mechanism can be compared where authorized. Public research can screen access and payout feasibility, but does not prove account eligibility or authorize acquisition.

Output and exit: one recorded route decision, important assumptions, the cheapest next uncertainty to resolve, and stop/revisit conditions. If access or data makes a route infeasible, stop or redesign it. The existing XAUUSD/M1 choice changes only with the required evidence and owner decision.^{1,3}

3. Research brief and experiment design

Why chosen: it makes the research answer a specific decision and supplies a stable specification that people and automation can reuse.

Define the hypothesis, comparator, evidence claim, fixed candidate version, permitted data partitions, cost assumptions, validity tests, acceptance rules and output format. Bound researcher time as well as compute time. Separate discovery from confirmation, and identify which new information would justify another iteration.³

Output and exit: an experiment brief and versioned specification, with an explicit pass, fail, inconclusive and invalid-evidence disposition. Confirm that the requested conclusion is achievable with the proposed data before expensive execution. Every required approval has an identified decision owner.

The sequence is conditional. Phases 1-3 can be brief when their inputs are already current. A small technical probe can legitimately resolve feasibility during Phase 2; “research before development” does not mean that every useful experiment must wait for exhaustive desk research.

1. py-mynyra, WORLDMAP.md.

3. py-mynyra, Locked research expansion sequence.

10. Project Management Institute, Projects and the Project Lifecycle.

12. Association for Project Management, What is benefits management and project success?.

Phases 4-6: produce evidence the next decision can use

4. Data, costs and evidence readiness

Why chosen: an experiment can be perfectly executed yet incapable of supporting its intended conclusion because the necessary observations are missing.

Reuse validated datasets and cost evidence where their scope still applies. Check provenance, timestamp and price-side semantics, gaps, coverage, independent periods and the relevance of cost samples. Freeze exact inputs and register their references in the existing catalog. New data ingestion deserves its own bounded work package only when a material gap requires it.⁵

Output and exit: a readiness decision for the named claim. In the inspected batch, March's available weekdays could not satisfy the registered 60-active-day requirement and longer independent history was unavailable. This was an acknowledged exploratory fallback. A useful preflight would expose that limitation before launch; it must not weaken the threshold or erase the separate economic failures.^{7,26}

5. Exploratory implementation and comparison

Why chosen: plausible strategy ideas need local, after-cost evidence, produced with less repeated setup and coordination work.

Implement only the behavior needed to test the fixed hypothesis. Reuse the simulator, reference controls, catalog and verified interfaces. Run bounded jobs, record progress and failures, retain partial attempts, and generate the same result package each time. Added variants are new recorded hypotheses; more cases are not automatically more independent evidence.^{3,7}

Output and exit: a complete, internally valid exploratory comparison and a clear disposition for every candidate. Measure operator handling and rework separately from machine runtime. A rejected strategy closes its work package; it does not automatically authorize another campaign.

6. Independent validation and advancement decision

Why chosen: selection on familiar data must not be mistaken for evidence that survives a genuinely independent test.

Use an unchanged candidate only after an authorized, frozen protocol and appropriate independent observations exist. Review model assumptions, cost sensitivity, stability and important accounting distinctions. A second deterministic execution establishes reproducibility; it does not create new market observations or independent economic validation.¹⁵

Output and exit: advance, reject, remain inconclusive, or invalidate, with evidence and authority recorded separately. In the present state there is no advancement candidate and the reserved period is sealed. This phase is blocked for advancement; the lifecycle does not create a reason to open April.⁷

The loop returns to the phase responsible for a gap: missing observations to Phase 4, a new research question to Phase 3, and a doubtful economic route to Phase 2. Corrected code or data creates a new evidence version with explicit lineage. Repeated use of the same observations remains adaptive research.

3. py-mynyra, Locked research expansion sequence.

5. py-mynyra, Praxis Step 3: one frozen data-to-result path.

7. py-mynyra, Praxis Step 4: expanded exploratory comparison.

15. Federal Reserve, OCC and FDIC, Supervisory Guidance on Model Risk Management.

Phases 7-9: connect evidence to operating benefit

7. Operational validation and controlled transition

Why chosen: historical evidence does not establish that the intended process can observe, execute, stop and recover reliably in its actual environment.

For the automated trading branch, first observe a qualified, unchanged candidate in the prescribed view-only shadow setting. A later demo-order experiment is a separate authorized transition with its own reconciliation, failure and recovery tests. A feed sample, shadow observation and demo execution are different evidence products.¹⁶

Output and exit: reviewed readiness evidence for the next explicitly bounded action, with runbook, observability and recovery responsibility. The earlier FIBO connection and spread sample remain useful inputs, but do not constitute this whole phase. No current strategy has reached it.⁷

8. Bounded economic pilot and cash realization

Why chosen: a working system and even account profit are insufficient if the project cannot receive usable cash.

Only after appropriate evidence, current account/access checks and exact owner authorization, evaluate a bounded economic attempt. Record the maximum committed capital, downside, failure residue and stop rules. Distinguish simulated results, booked profit, withdrawal eligibility, requested withdrawal and received cash. Recheck the selected route's actual conditions before relying on them.¹

Output and exit: reconciled evidence of the pilot's costs, outcome and any cash actually received, or a recorded stop and recoverable remaining state. One receipt demonstrates a receipt, not repeatable self-financing operation. This report proposes no provider, position, purchase or trading action.

9. Repeatability, operating handover and benefits review

Why chosen: the destination is continuing financial support for operations, not completion of the bot or one profitable episode.

Assess multiple observed periods and actual receipts against recurring operating costs, obligations and required reserves. Confirm that normal failures can be handled without an unsustainable increase in operator effort. Define who owns routine operation and when degraded evidence sends the activity back for review. Benefits need explicit ownership beyond delivery of a project artifact.¹²

Output and exit: evidence that the agreed self-sustaining threshold is met over the agreed window, or a decision that it remains unproven. Close the establishment project only against that definition; ongoing monitoring and subsequent improvements remain operations or separately scoped projects.

Phases 7-9 describe conditional future transitions. They are not a mandate to complete every technical layer. A route chosen in Phase 2 determines the necessary operational evidence. Across all nine phases, maintain decisions, scope changes, costs, risks, evidence links and knowledge transfer as continuous activities.

1. py-mynyra, WORLDMAP.md.

7. py-mynyra, Praxis Step 4: expanded exploratory comparison.

12. Association for Project Management, What is benefits management and project success?.

16. Financial Conduct Authority, Multi-firm review of algorithmic trading controls: high-level observations.

A research subproject that leaves a reusable result

Strategy research spans Phases 3-6 and can be managed as one bounded subproject when its scope warrants it. Data acquisition/normalization or operational validation can likewise become separate subprojects if they have distinct deliverables and acceptance. A heavy task alone does not require a new management hierarchy.

Subphase	Deliverable and decision
A. Intake and reuse check	Name the decision; load the latest relevant experiment, rejected approaches, approved scope and reusable inputs. Identify exactly what is new.
B. Source and hypothesis research	A short evidence map connecting each claim to its source, applicability, uncertainty and proposed test. Stop browsing when more sources are unlikely to change the next decision.
C. Readiness assessment	A data/cost/coverage assessment that states whether the experiment can support exploration or qualification. Route missing inputs to their owner.
D. Specification and contract proof	Version the protocol, inputs and outputs. Check baseline compatibility and result identity before the expensive batch.
E. Bounded execution	Produce complete results through existing deterministic scripts; record interruptions and preserve invalid evidence explicitly.
F. Review and decision	Reproduce the numerical result where required; assess the claim at the correct evidence level; record advancement or no-go authority separately.
G. Closeout and next-run handoff	Publish an appropriate sanitized summary, retained evidence references, lessons, unresolved blockers and next admissible action. Reuse it at intake.

The research progression is adapted from CRISP-DM's business understanding, data understanding, data preparation, modeling, evaluation and deployment stages. Here "deployment" of a research result means putting a reviewed decision and reusable knowledge into use; it does not mean deploying a trading strategy.^{13,28}

Definition of done for an unsuccessful campaign: the result is valid at its stated scope; the decision is recorded; reusable data, tests and workflows are linked; the next operator can find the state; and any repeated friction has a disposition. This allows a campaign to fail its strategy objective while still producing reusable capability.

A lesson should change an input, contract, test or workflow where justified. For example, the observed experiment-identity collision suggests a cross-experiment artifact compatibility check before future batches. That is a focused follow-through from a known defect, not a reason to redesign the whole catalog.⁶

Repeated completion without a usable handoff is a process problem even when the numerical work is rigorous. Conversely, re-reading familiar data or rerunning accounting checks can remain necessary; reducing friction must preserve the meaning and provenance of the evidence.

6. py-mynyra, Step 4 artifact identity correction.

13. IBM, CRISP-DM View, SPSS Modeler 18.6.0.

25. Association for Project Management, Project management glossary.

28. IBM, Understanding and preparing data.

Shared project state and communication contracts

The minimal missing capability to test is a reliable answer to: **What is current, why is it current, who owns the next decision, and what may happen next?** The answer can be assembled from existing authoritative records. It does not require one database to own every fact.

Information	Authority and handoff
Goals, scope and decisions	Reviewed project records own these facts. Each status projection cites their version and decision date.
Dataset and experiment identity	Existing catalog relations and immutable manifests own provenance and lineage. Keep the current division between metadata and evidence files.
Task execution state	A runner reports pending, running, failed or complete with a unique execution ID and output reference. Completion is not an advancement decision.
Current project view	A generated summary joins approved decisions, latest valid evidence, unresolved blockers, next owner and last-refresh time. It must expose stale or unavailable inputs.
AI context	An AI reads the same selected records with version and access boundaries. Its narrative is a derived view; it cannot silently become project authority.

For each handoff, define a compact contract: project/work-package ID, experiment ID, execution ID, code/config/data versions, evidence references, validity state, decision status, responsible owner and observed-at time. Use existing identifiers where possible. The catalog already supplies much of the evidence identity; the operating layer supplies the remaining coordination facts.⁸

A proposed completion path is: the runner records immutable outputs; verification checks completion and integrity; a summary is generated; a reviewer records the conclusion; the approved decision updates the current view. A failure path preserves the attempt, identifies its owner and proposes a bounded repair. A retry gets a new execution identity linked to the prior attempt.

Some events can trigger the next already-authorized mechanical step. Material decisions still need the existing owner’s authority. An approval belongs to its exact scope and versions; it is not a reusable permission token for later datasets, strategies or exposure.

BCBS guidance emphasizes authoritative definitions, reliable aggregation and traceable data ownership. The transferable point is coherence across records, not the choice of a particular database engine.¹⁷ MCP may make these records accessible to AI clients, but it cannot decide which conflicting record is authoritative.²⁰

Keep private evidence on the authorized host and share only permitted references or sanitized summaries. Ordinary code synchronization does not synchronize ignored evidence. A shared summary must accurately state when it cannot access the evidence it references.

8. py-mynyra, Authoritative local SQLite catalog migration.
17. Basel Committee on Banking Supervision, Risk data aggregation and risk reporting, SRP36.
20. Model Context Protocol, Architecture overview.

Tool choices and the economics of automation

The initial recommendation is to reuse the present repository, immutable artifacts, SQLite catalog and deterministic runners. Add a small status-and-handoff surface only around a demonstrated gap. Evaluate larger tools against saved operator effort and new maintenance burden.

Candidate	Role and adoption condition
Existing Git records and task view	Use a linked phase/work-package index and a current status projection. Prefer an existing shared task surface if it already meets the need.
SQLite or PostgreSQL	SQLite remains consistent with the measured local choice. Consider a service/API or PostgreSQL if real remote access, independent identities or concurrent writers require it; do not sync a live database file as a coordination shortcut. ^{4,21}
Python runner; n8n if needed	Keep numerical execution deterministic. A workflow tool can route events and failures between services when that removes repeated handoffs. n8n documents execution-error workflows. ²²
MCP interface	Useful when AI clients repeatedly need access to the same allowed state and evidence. It is an access interface, not shared memory, a task queue or a synchronization policy. ²⁰
Langflow	Useful for a repeatable AI research flow assembled from components. Its visual workflow editor does not establish evidence quality or authorize the model's conclusions. ²³
MLflow, only after a concrete comparison need	Provides experiment/run tracking of parameters, metrics and artifacts. Evaluate whether its interface adds enough value over the existing catalog to justify another metadata surface. ²⁴

SQLite supports one writer at a time, but an application service can serialize brief writes and serve remote readers. Multiple people do not automatically require PostgreSQL. Direct network-filesystem sharing and copying an active SQLite file are different designs from a controlled service or consistent snapshot.²¹

An illustrative time-payback check: a connector takes 6 hours to build and maintain during the decision horizon and saves 30 minutes on each comparable run. It needs 12 such runs to recover those hours, before considering failures or cash expense. These are example values, not measurements of Mynyra. A smaller change that saves the same handling time would be preferable.

Automation candidates include generating evidence indexes, validating handoff fields, assembling status summaries and routing known failures. Research judgment, acceptance of weak evidence and material economic choices cannot be reduced to a success exit code. Existing project authority determines which automatic transitions are allowed.

No new tools, services, schedules or external messages have been configured by this assessment. The recommendations are deliberately conditional on observing the actual cross-machine workflow.

4. py-mynyra, Praxis Step 2: catalog and storage benchmark.
20. Model Context Protocol, Architecture overview.
21. SQLite, Appropriate Uses For SQLite.
22. n8n, Error handling.
23. Langflow, Use the visual editor.
24. MLflow, MLflow Tracking.

Verification of the proposed management improvement

The next useful management experiment is a bounded review of one representative handoff, followed by one small repair and a comparable measurement. It need not run another strategy campaign. Use existing permitted evidence to test whether another operator can identify the current result and the next admissible action.

Measure	What it resolves
Time to recover current context	Minutes needed to locate the right decision, experiment, dataset and unresolved blocker, without a narrated reconstruction.
Operator touch time	Time spent issuing routine commands, copying references, fixing paths and manually assembling summaries. Record setup, novelty and judgment separately.
Handoff wait and clarification	Elapsed time between evidence availability and usable receipt by the next person; count missing information that causes a return.
Rework and invalid attempts	Extra effort caused by identity, input, environment or workflow defects. Distinguish this from an economically rejected candidate.
Reuse and maintenance	Which existing assets carry forward unchanged, and how much continuing work the new integration adds.

DORA's value-stream approach is useful for mapping actual work and delays before changing tools.¹⁹ Google's toil concept helps distinguish repetitive operational handling from novel engineering. Its definition is about production operations; applying it to research is an analogy, and two campaigns alone do not quantify a stable pattern.¹⁸

Proposed acceptance: a second operator can recover the reviewed state and evidence links without a custom briefing; a representative handoff requires less measured handling; partial or stale output cannot appear approved; and no scientific check or authority boundary is bypassed. Set the time target from the baseline rather than invent an unsupported percentage improvement.

Incremental adoption: first link the nine phases to existing records and identify current blockers; then measure a representative handoff; then implement only the smallest justified summary or contract fix. Consider remote services only if this reveals an actual access or concurrency limitation. These are recommended next work packages, not changes already performed.

For each proposed repair, record its economic movement (operator time returned to useful work), information value (the bottleneck it tests), irreversible downside (build time and any expense), and failure residue (the original workflow and retained evidence). A small reversible integration can be evaluated on those terms without committing to a new platform.

Remaining uncertainty: the amount of coordination time, current private tools, actual host access, operator responsibilities and economic thresholds are not verified here. The second batch had twice as many registered cases plus catalog/recovery work; similar overall difficulty therefore does not by itself show that reuse failed.⁷ The evidence supports targeted integration work, not a claim that project management is the sole cause of the six-month outcome.

7. py-mynrya, Praxis Step 4: expanded exploratory comparison.

18. Google SRE, Vivek Rau, Eliminating Toil.

19. DORA, Andrew Davis and Rob Edwards, How to use value stream mapping to improve software delivery.

Sources: project records and management foundations

Repository file links are pinned to the inspected commit. They establish what the source and documentation contain, not independent reproduction of private run claims. PR descriptions may change after publication. Official explanatory pages are used for concepts; no claim of PMI certification or regulatory applicability is made.

1. **py-mynyra.** WORLDMAP.md Status 3 Sep; progress 6 Sep 2026.
2. **py-mynyra.** IMPLEMENTATION_PLAN.md and README.md Plan status 6 Sep 2026; README inspected at the same commit.
3. **py-mynyra.** Locked research expansion sequence Agreed 7 Sep 2026; includes Step 4 closeout.
4. **py-mynyra.** Praxis Step 2: catalog and storage benchmark Measured 7 Sep 2026.
5. **py-mynyra.** Praxis Step 3: one frozen data-to-result path Completed evidence 7 Sep 2026.
6. **py-mynyra.** Step 4 artifact identity correction Inspected 11 Sep 2026; dated via the referenced commits.
7. **py-mynyra.** Praxis Step 4: expanded exploratory comparison Completed work package; inspected 11 Sep 2026.
8. **py-mynyra.** Authoritative local SQLite catalog migration Inspected 11 Sep 2026.
9. **py-mynyra.** Praxis and bounded Step 4 runner Read with scripts/praxis-step4.py at the same commit.
10. **Project Management Institute.** Projects and the Project Lifecycle Undated current explanatory page; accessed 11 Sep 2026.
11. **Association for Project Management.** What is a life cycle? APM Body of Knowledge 8th-edition definition; accessed 11 Sep 2026.
12. **Association for Project Management.** What is benefits management and project success? Current explanatory page; accessed 11 Sep 2026.
13. **IBM.** CRISP-DM View, SPSS Modeler 18.6.0 Versioned official documentation; accessed 11 Sep 2026.
25. **Association for Project Management.** Project management glossary Work breakdown structure and work package entries; accessed 11 Sep 2026.
26. **py-mynyra.** PR #12: Strategy Selection Research: Step 4 qualification and no-go decision Merged 11 Sep 2026; PR description is mutable.
27. **Mynyra-Trade.** Mynyra_Demo_Phase_Charter_v1.docx Existing private record, 7 pages; stored 30 Aug 2026. Related-project context, not an adopted py-mynyra requirement.

Additional inspected paths: README.md; AGENTS.md; docs/COMPARISON_WORKFLOW.md; docs/STRATEGY_RESEARCH_PLAN.md; docs/PRAXIS_STEP4_PROTOCOL.md; src/mynyra/catalog.py; scripts/praxis-step3.py; scripts/praxis-step4.py. All at the pinned repository commit.

Sources: financial controls and operating workflows

Financial-sector publications are applied proportionally as design references. The current revised US model-risk guidance is used in place of superseded SR 11-7. Mutable software documentation was checked for present conceptual capabilities; no installation, compatibility or service-availability guarantee is inferred.

14. Federal Reserve, OCC and FDIC. SR 26-2: Revised Guidance on Model Risk Management 17 Apr 2026; supersedes SR 11-7 and SR 21-8.

15. Federal Reserve, OCC and FDIC. Supervisory Guidance on Model Risk Management Revised 2026 guidance; sections IV-VI.

16. Financial Conduct Authority. Multi-firm review of algorithmic trading controls: high-level observations Published and updated 21 Aug 2025.

17. Basel Committee on Banking Supervision. Risk data aggregation and risk reporting, SRP36 15 Dec 2019 framework version; particularly paragraphs 36.18 and 36.27.

18. Google SRE, Vivek Rau. Eliminating Toil Site Reliability Engineering, 2016, chapter 5; accessed 11 Sep 2026.

19. DORA, Andrew Davis and Rob Edwards. How to use value stream mapping to improve software delivery Current guide; accessed 11 Sep 2026.

20. Model Context Protocol. Architecture overview Current conceptual documentation; accessed 11 Sep 2026.

21. SQLite. Appropriate Uses For SQLite Current documentation; accessed 11 Sep 2026.

22. n8n. Error handling Current documentation; accessed 11 Sep 2026.

23. Langflow. Use the visual editor Current documentation; accessed 11 Sep 2026.

24. MLflow. MLflow Tracking Current documentation; accessed 11 Sep 2026.

28. IBM. Understanding and preparing data Official indexed documentation; direct page retrieval returned 403. Used only for its listed CRISP-DM stages.

Assessment date: 11 September 2026. Scope: py-mynyra project management and research-workflow design. The nine-phase sequence, subproject design, operating contracts and proposed measurements are analytical recommendations synthesized for this project; they are not a published universal financial-project lifecycle.